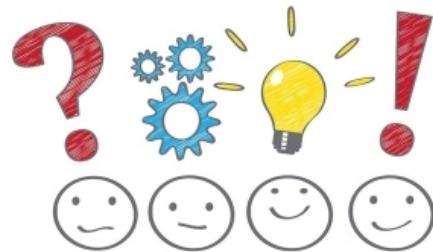


# Admin

*First lab and assign*

We're off and running!

*Check in*



## Today: From Assembly to C (and back again)

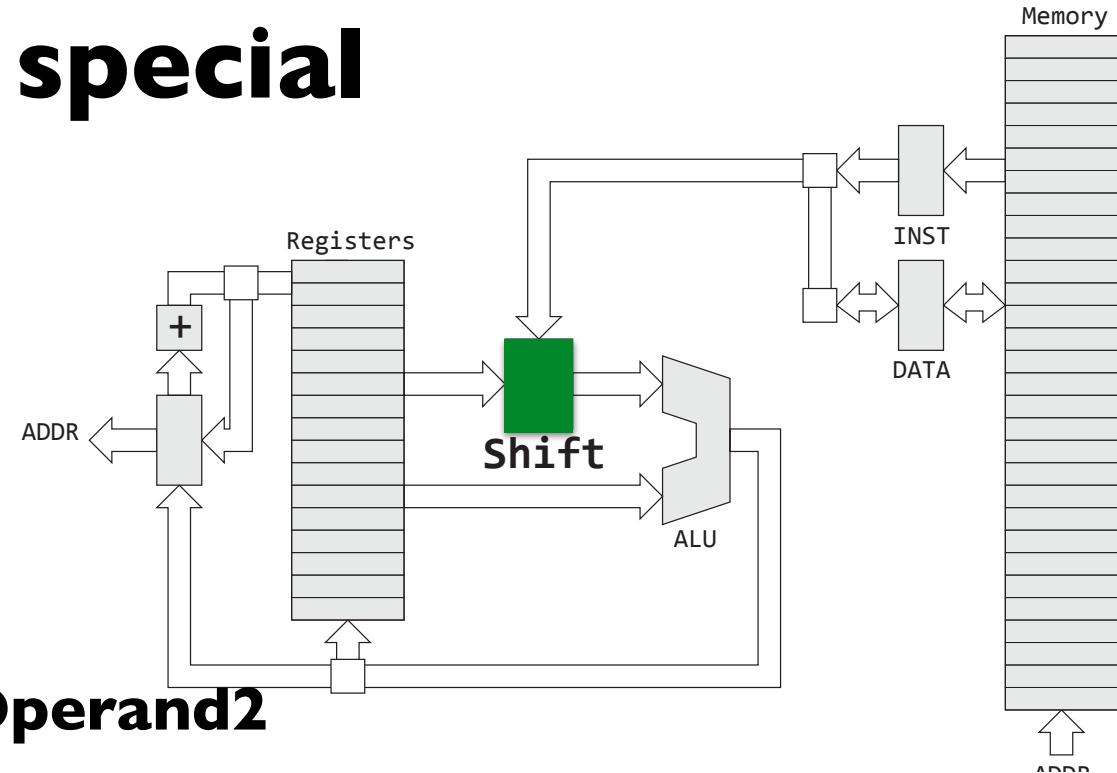
ARM condition codes, branch instructions

C language as “high-level” assembly

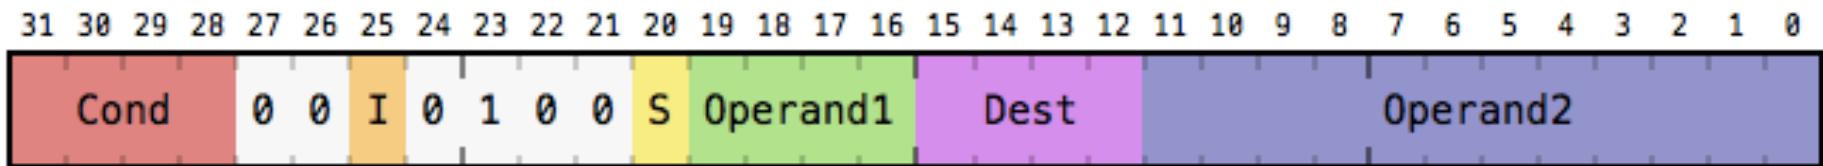
What does a compiler do?

Makefiles

# Operand 2 is special



**Dest = Operand1 op Operand2**



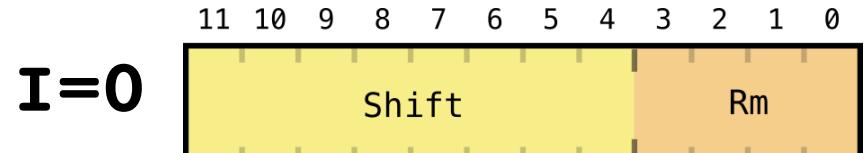
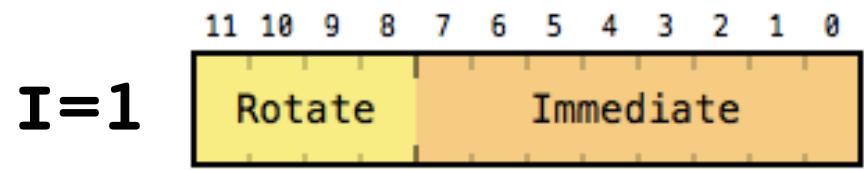
**add r0, r1, #0x1f000**

**sub r0, r1, #6**

**rsb r0, r1, #6**

**add r0, r1, r2, lsl #3**

**mov r1, r2, ror #7**

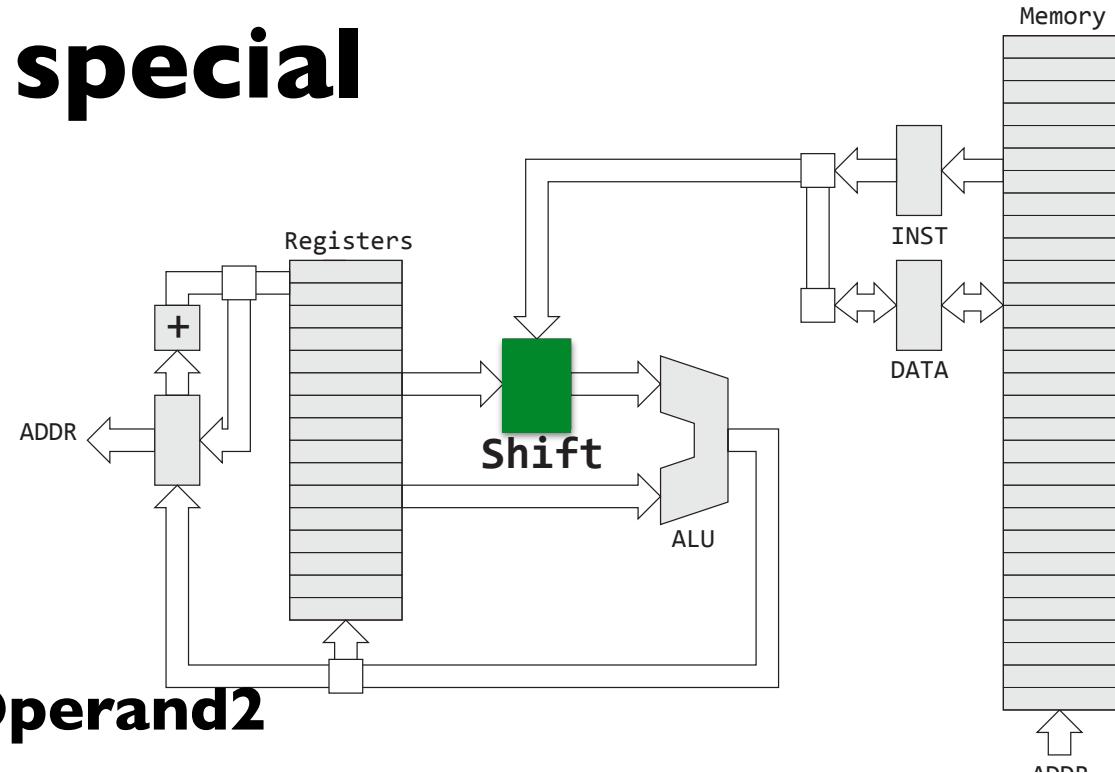


```
/// SET0 = 0x2020001c
mov r0, #0x20          // r0 = 0x00000020
lsl r1, r0, #24        // r1 = 0x20000000
lsl r2, r0, #16        // r2 = 0x00200000
orr r0, r1, r2         // r0 = 0x20200000
orr r0, r0, #0x1c      // r0 = 0x2020001c
```

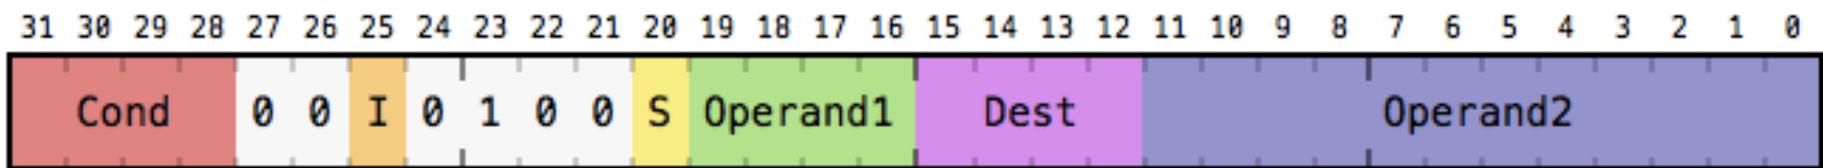
```
// SET0 = 0x2020001c
mov r0, #0x20000000 // 0x20>>>8
orr r0, #0x00200000 // 0x20>>>16
orr r0, #0x0000001c // 0x1c>>>0
```

Rotation of small constants makes code  
40% shorter (and 40% faster)

# Operand 2 is special



**Dest = Operand1 op Operand2**



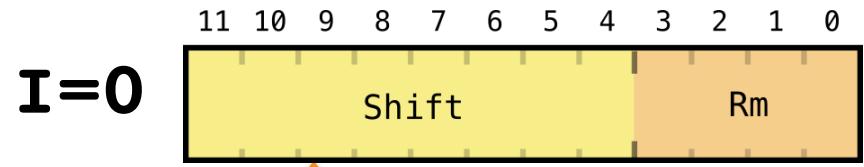
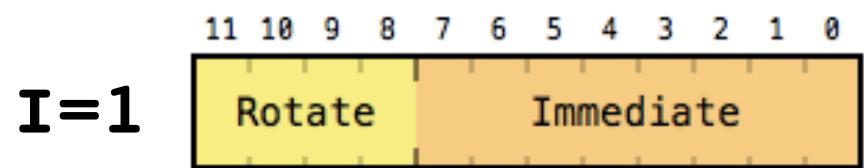
**add r0, r1, #0x1f000**

**sub r0, r1, #6**

**rsb r0, r1, #6**

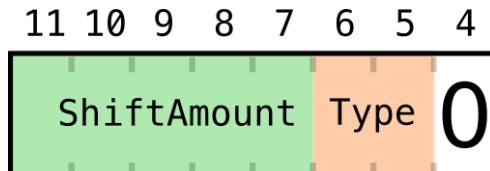
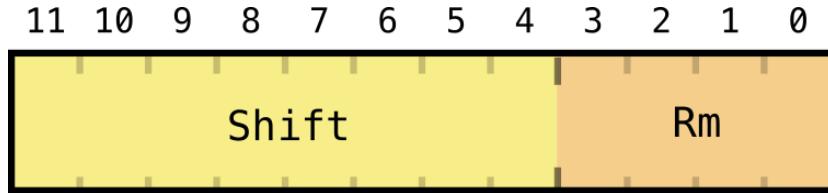
**add r0, r1, r2, lsl #3**

**mov r1, r2, ror #7**

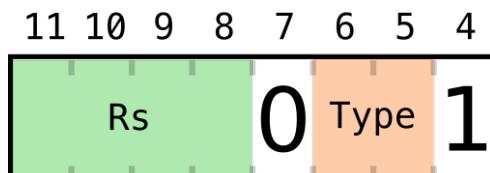


**lsl, lsr, asr, ror**

# Shift Operand Details



**Shift Rm by ShiftAmount**



**Shift Rm by Rs**

**00: lsl (logical shift left)**

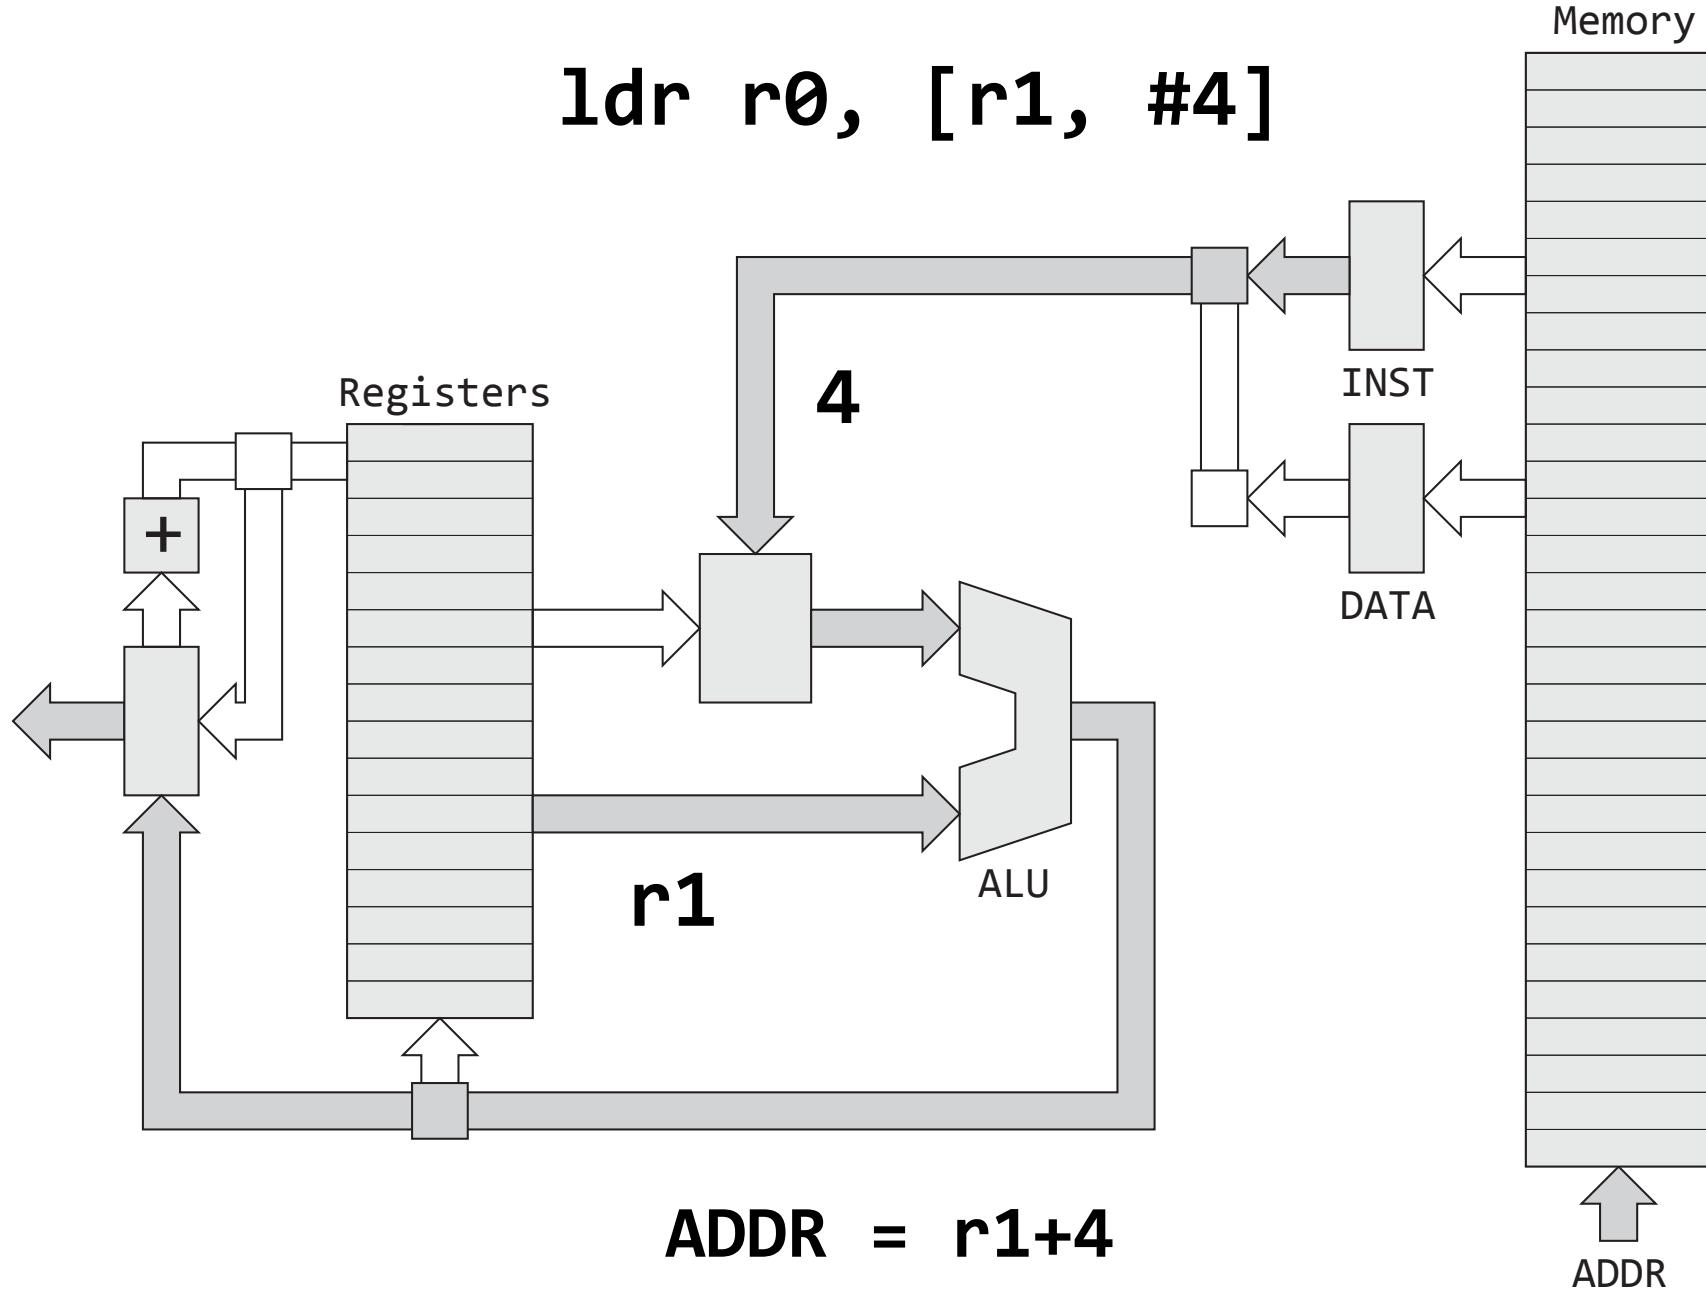
**01: lsr (logical shift right)**

**10: asr (arithmetic shift right)**

**11: ror (rotate right)**

# Load from Memory to Register (LDR)

**ldr r0, [r1, #4]**



```
// configure GPIO 20 for output  
ldr r0, =0x20200008  
mov r1, #1  
str r1, [r0]
```

```
// set bit 20  
ldr r0, =0x2020001C  
mov r1, #0x00100000  
str r1, [r0]
```

```
loop: b loop
```

```
// configure GPIO 20 for output
ldr r0, [pc + 20]
mov r1, #1
str r1, [r0]
```

```
// set bit 20
ldr r0, [pc + 12]
mov r1, #0x00100000
str r1, [r0]
```

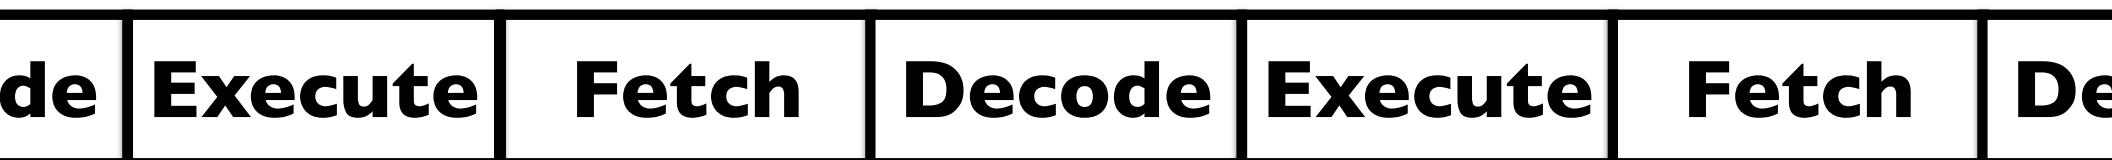
loop: b loop

```
.word 0x20200008
.word 0x2020001C
```

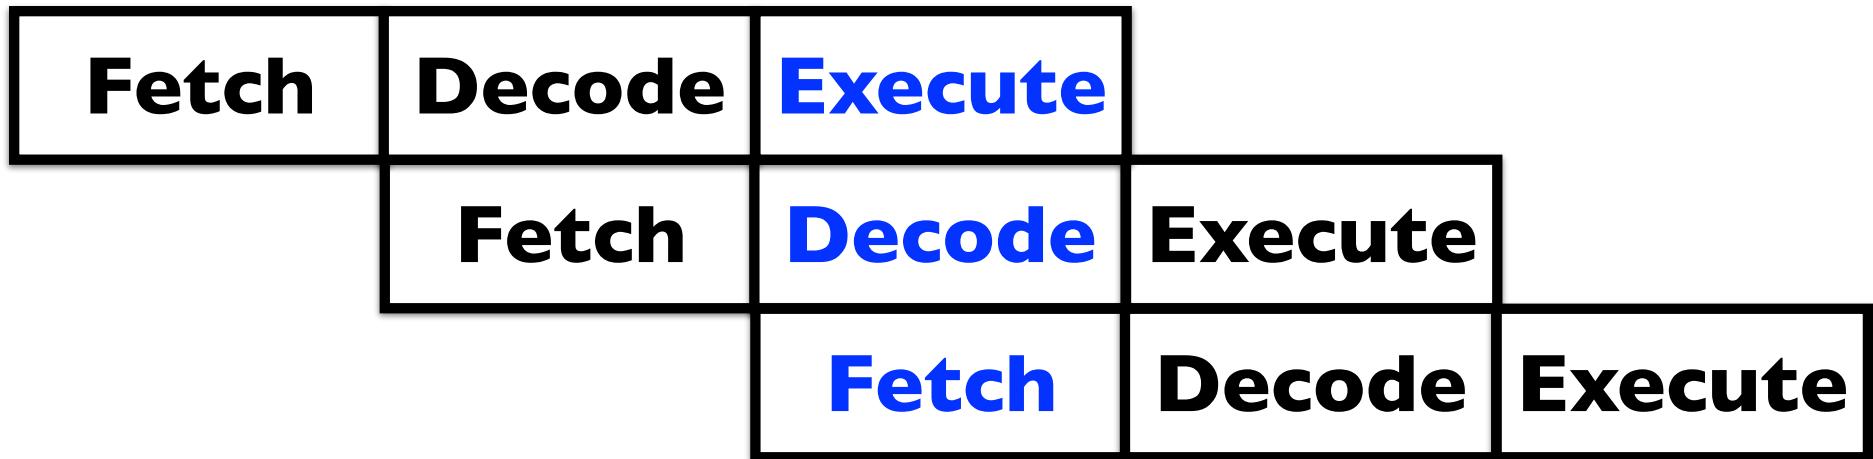
# **3 steps to run an instruction**

**Fetch    Decode    Execute**

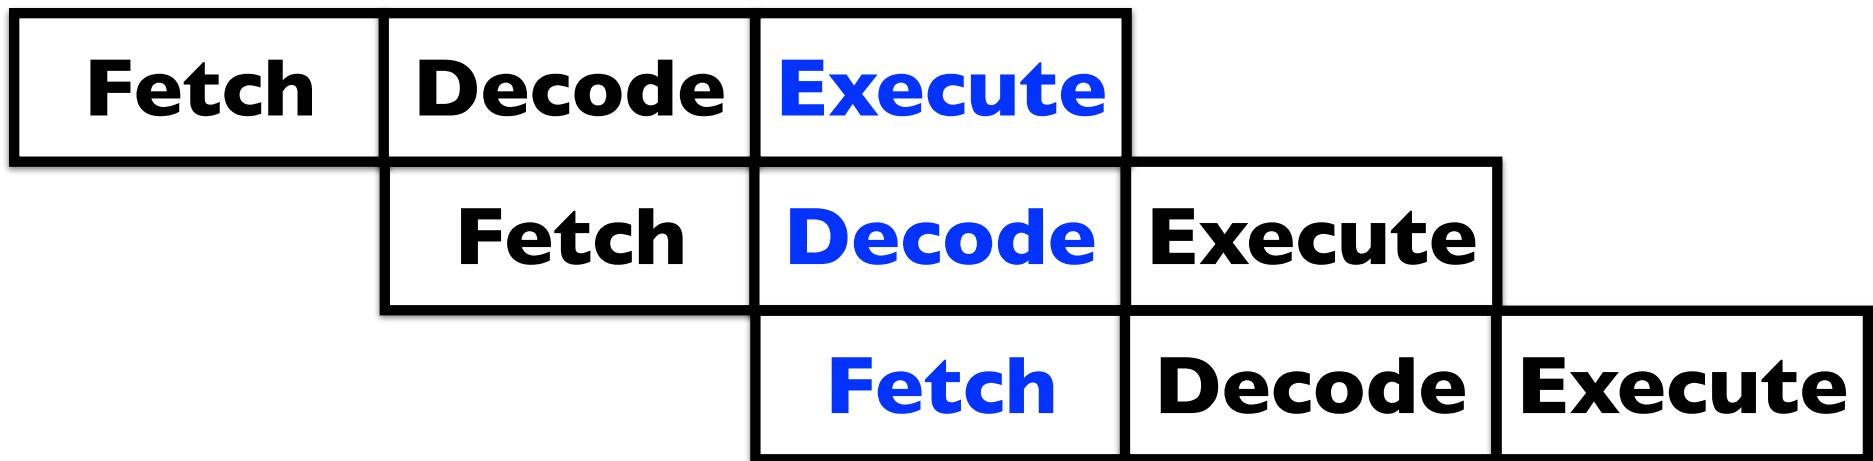
# **3 instructions takes 9 steps**



**To speed things up,  
steps are overlapped ("pipelined")**



**To speed things up,  
steps are overlapped ("pipelined")**



**PC value in the executing instruction is equal to  
the pc value of the instruction being fetched -  
which is 2 instructions ahead (PC+8)**

# **Changing Control Flow**

## **Loops, branches, and condition codes**

# Control flow

Instructions stored in memory contiguously

Register **pc** (**r15**) tracks address in memory where instructions are being read

Default is "straight-line" code: next instruction to execute is at next word address (**pc = pc + 4**)

**branch** instructions change what instruction is fetched, decoded, and executed next has effect of **pc = target**

**b target**

Above branch is unconditional (always taken)

Branches can also be predicated on state of "condition codes"

# Condition Codes

**Z** result was 0

**N** result was < 0

**C** operation generated carry

**V** operation had arithmetic overflow

*(More on carry and overflow in later lecture...)*

## Which instructions set/clear codes?

**cmp** (like **sub**, but discards result)

**tst** (like **and**, but discards result)

Any data processing instruction suffixed with **s:**

**adds** **movs** **orrs** **lsrs** ...

**s bit**  
**(if on, instr will set condition codes)**



# Branch instructions

**b target**

**bne target**

**bmi target**

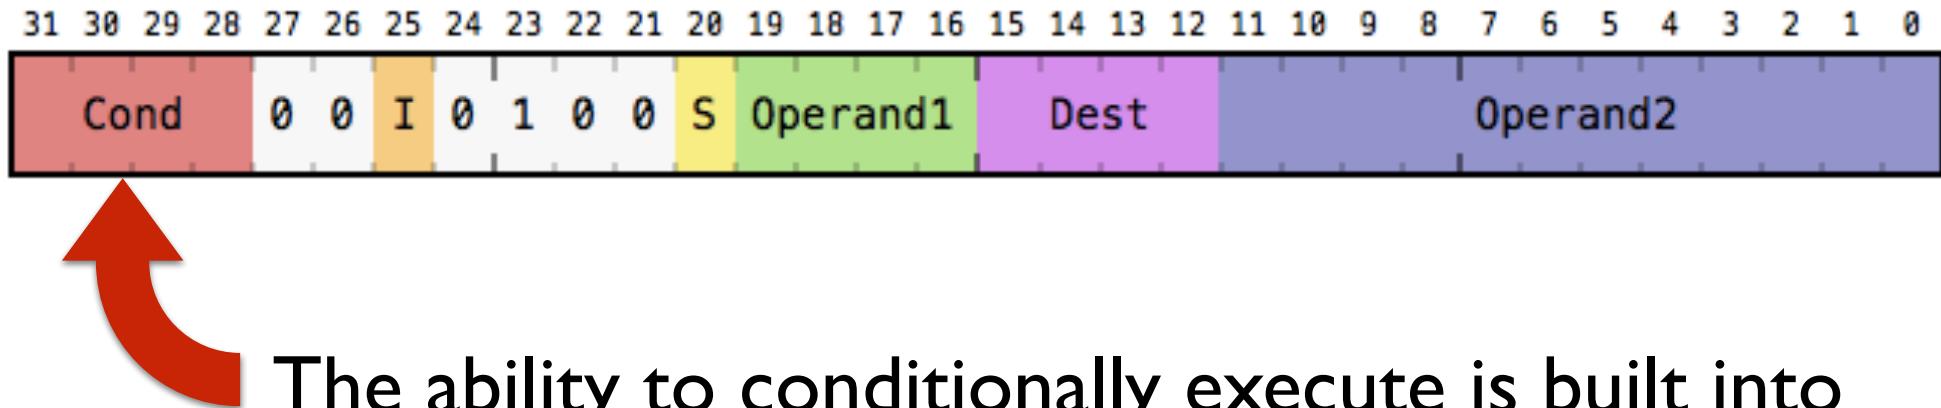
**bge target**

...

**branch** reads condition codes, as set by a previous instruction

If specific condition satisfied, branch is taken, **pc = target**  
otherwise falls through, **pc = pc + 4**

# Condition execution



The ability to conditionally execute is built into all ARM instructions! (not just **branch**...)

**addeq r0, r0, #3**  
**submi r1, r2, r3**

Q: Given our earlier foray into machine-encoded instructions, what do you suspect is the condition represented by **0xe**?

<b>Code</b>	<b>Suffix</b>	<b>Description</b>	<b>Flags</b>
0000	EQ	Equal / equals zero	Z
0001	NE	Not equal	!Z
0010	CS / HS	Carry set / unsigned higher or same	C
0011	CC / LO	Carry clear / unsigned lower	!C
0100	MI	Minus / negative	N
0101	PL	Plus / positive or zero	!N
0110	VS	Overflow	V
0111	VC	No overflow	!V
1000	HI	Unsigned higher	C and !Z
1001	LS	Unsigned lower or same	!C or Z
1010	GE	Signed greater than or equal	N == V
1011	LT	Signed less than	N != V
1100	GT	Signed greater than	!Z and (N == V)
1101	LE	Signed less than or equal	Z or (N != V)
1110	AL	Always (default)	any

# **Challenge for you all:**

**Write an assembly program to count the "on" bits in a given numeric value**

```
mov r0, #val  
mov r1, #0
```

```
// r0 holds input value  
// use r1 to store count of on bits in value
```

**<https://salmanarif.bitbucket.io/visual/>**

# VisUAL ARM Emulator

The screenshot shows the VisUAL ARM Emulator interface. At the top, there is a toolbar with buttons for New, Open, Save, Settings, Tools, Emulation Complete (with Line Issues 9 0), Execute, Reset, Step Backwards, and Step Forwards.

The main area displays assembly code:

```
1      mov    r0, #0x3a
2      mov    r1, #0
3
4
5      loop
6      tst    r0, #1
7      addne r1, r1, #1
8      lsrs   r0, r0, #1
9      bne    loop
10
```

A message "Reset to continue editing code" is displayed above the assembly code.

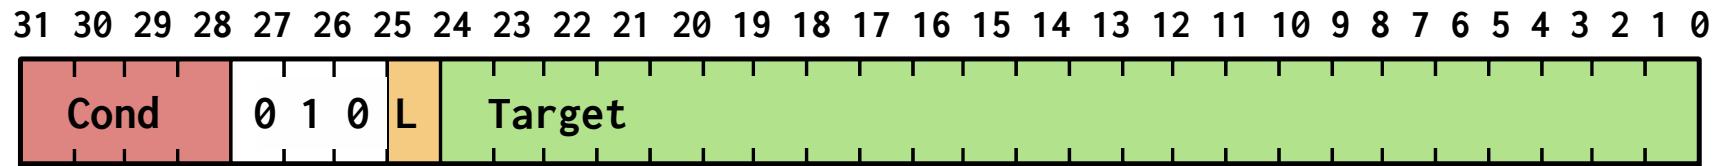
To the right, a register table shows the state of registers R0 through R12:

R0	0x0	Dec	Bin	Hex
R1	0x4	Dec	Bin	Hex
R2	0x0	Dec	Bin	Hex
R3	0x0	Dec	Bin	Hex
R4	0x0	Dec	Bin	Hex
R5	0x0	Dec	Bin	Hex
R6	0x0	Dec	Bin	Hex
R7	0x0	Dec	Bin	Hex
R8	0x0	Dec	Bin	Hex
R9	0x0	Dec	Bin	Hex
R10	0x0	Dec	Bin	Hex
R11	0x0	Dec	Bin	Hex
R12	0x0	Dec	Bin	Hex

At the bottom, there are status indicators for Clock Cycles (1 Total: 36) and CSPR Status Bits (NZCV) with values 0 1 1 0.

<https://salmanarif.bitbucket.io/visual/>

# Branch instruction encoding



**b (bal) branch always**

**1110 1010 tttt tttt tttt tttt tttt tttt**

**beq branch if zero CC set**

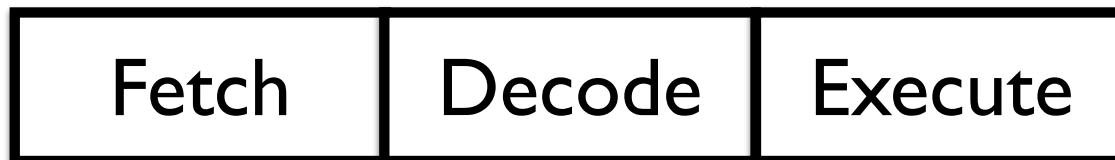
**0000 1010 tttt tttt tttt tttt tttt tttt**

branch target is PC-relative offset

green bits encode offset, counted in 4-byte words

*Q: How far can this reach?*

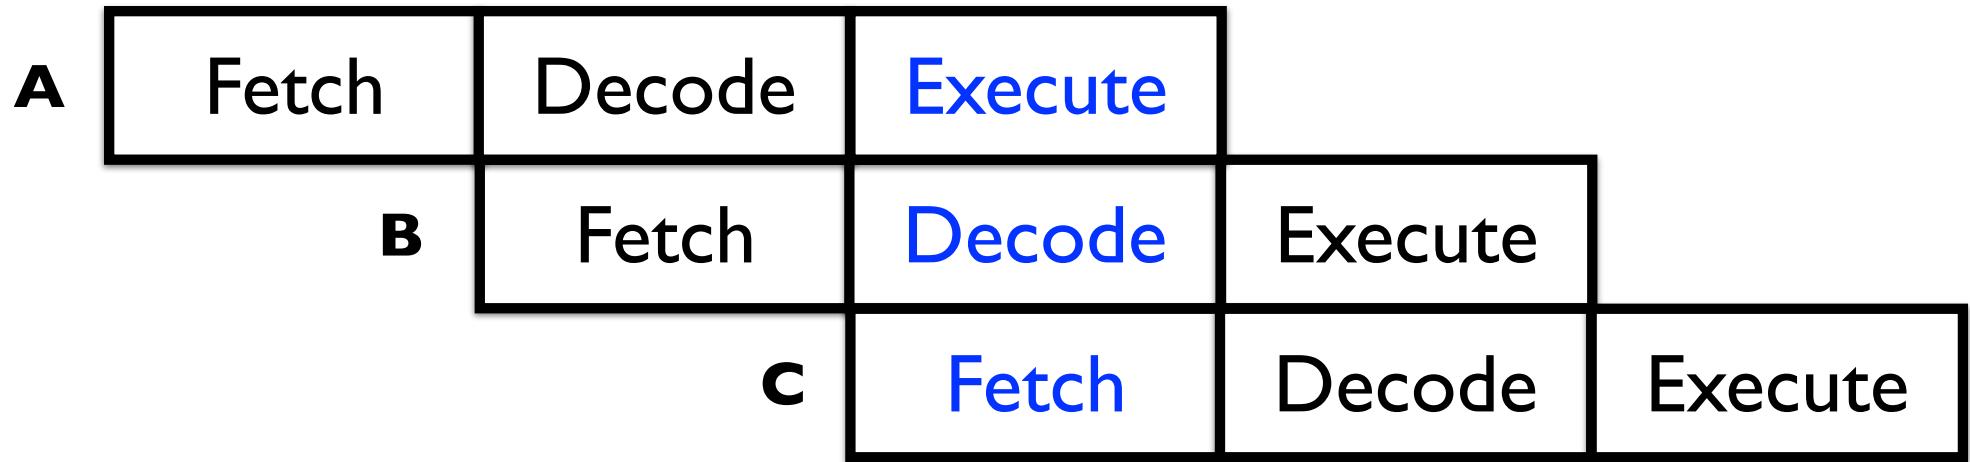
3 steps per instruction



3 instructions takes 9 steps in sequence

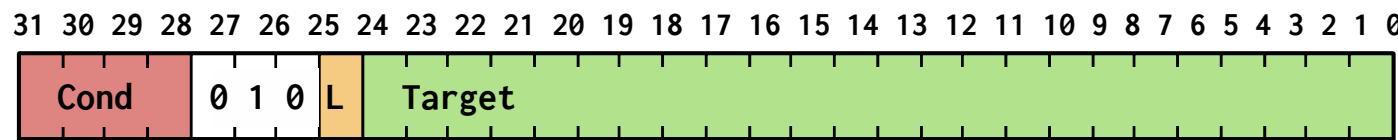


**To speed things up,  
steps are overlapped ("pipelined")**



During cycle that instruction A is executing, PC has advanced twice, holds address of instruction being fetched (C). This is 2 instructions past A (PC+8)

18: e3130001	tst r3, #1
1c: 1a000003	bne 30
20: e28300dc	add r0, r3, #107
24: e2433004	sub r3, r3, #4
28: e0800003	add r0, r0, r3
2c: e1a00003	mov r0, r3
30: e2833001	add r3, r3, #1



```
// 1a      branch if not equal
// 000003  target, expressed PC-relative as
//           count of 4-byte words
// this offset is 8 + 3*4 = 20 bytes ahead
```

# **ISA design is an art form!**

Some neat things about ARM design

Commonalities across operations

Register vs. immediate operands

Use of barrel shifter

All registers treated same (with a few caveats)

Predicated execution

Set condition code (or not)

Orthogonality leads to composability

# Why assembly?

What you see is what you get

No surprises

Precise control, timing

Unfettered access to hardware

# Why C?

More concise

Easier to read

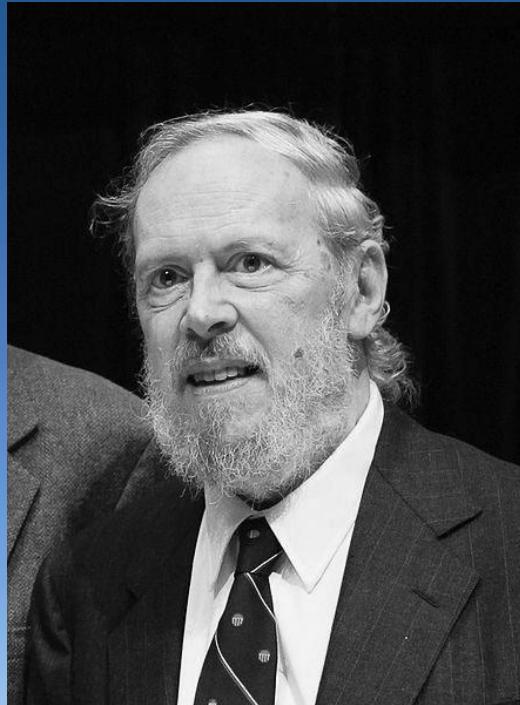
Can name variables and structures

Type-checking

More portable (ARMv6, ARMv7, ARMv8...)

Functions!!

*Real question is not whether to use assembly, but when...*



**Dennis Ritchie**

SECOND EDITION

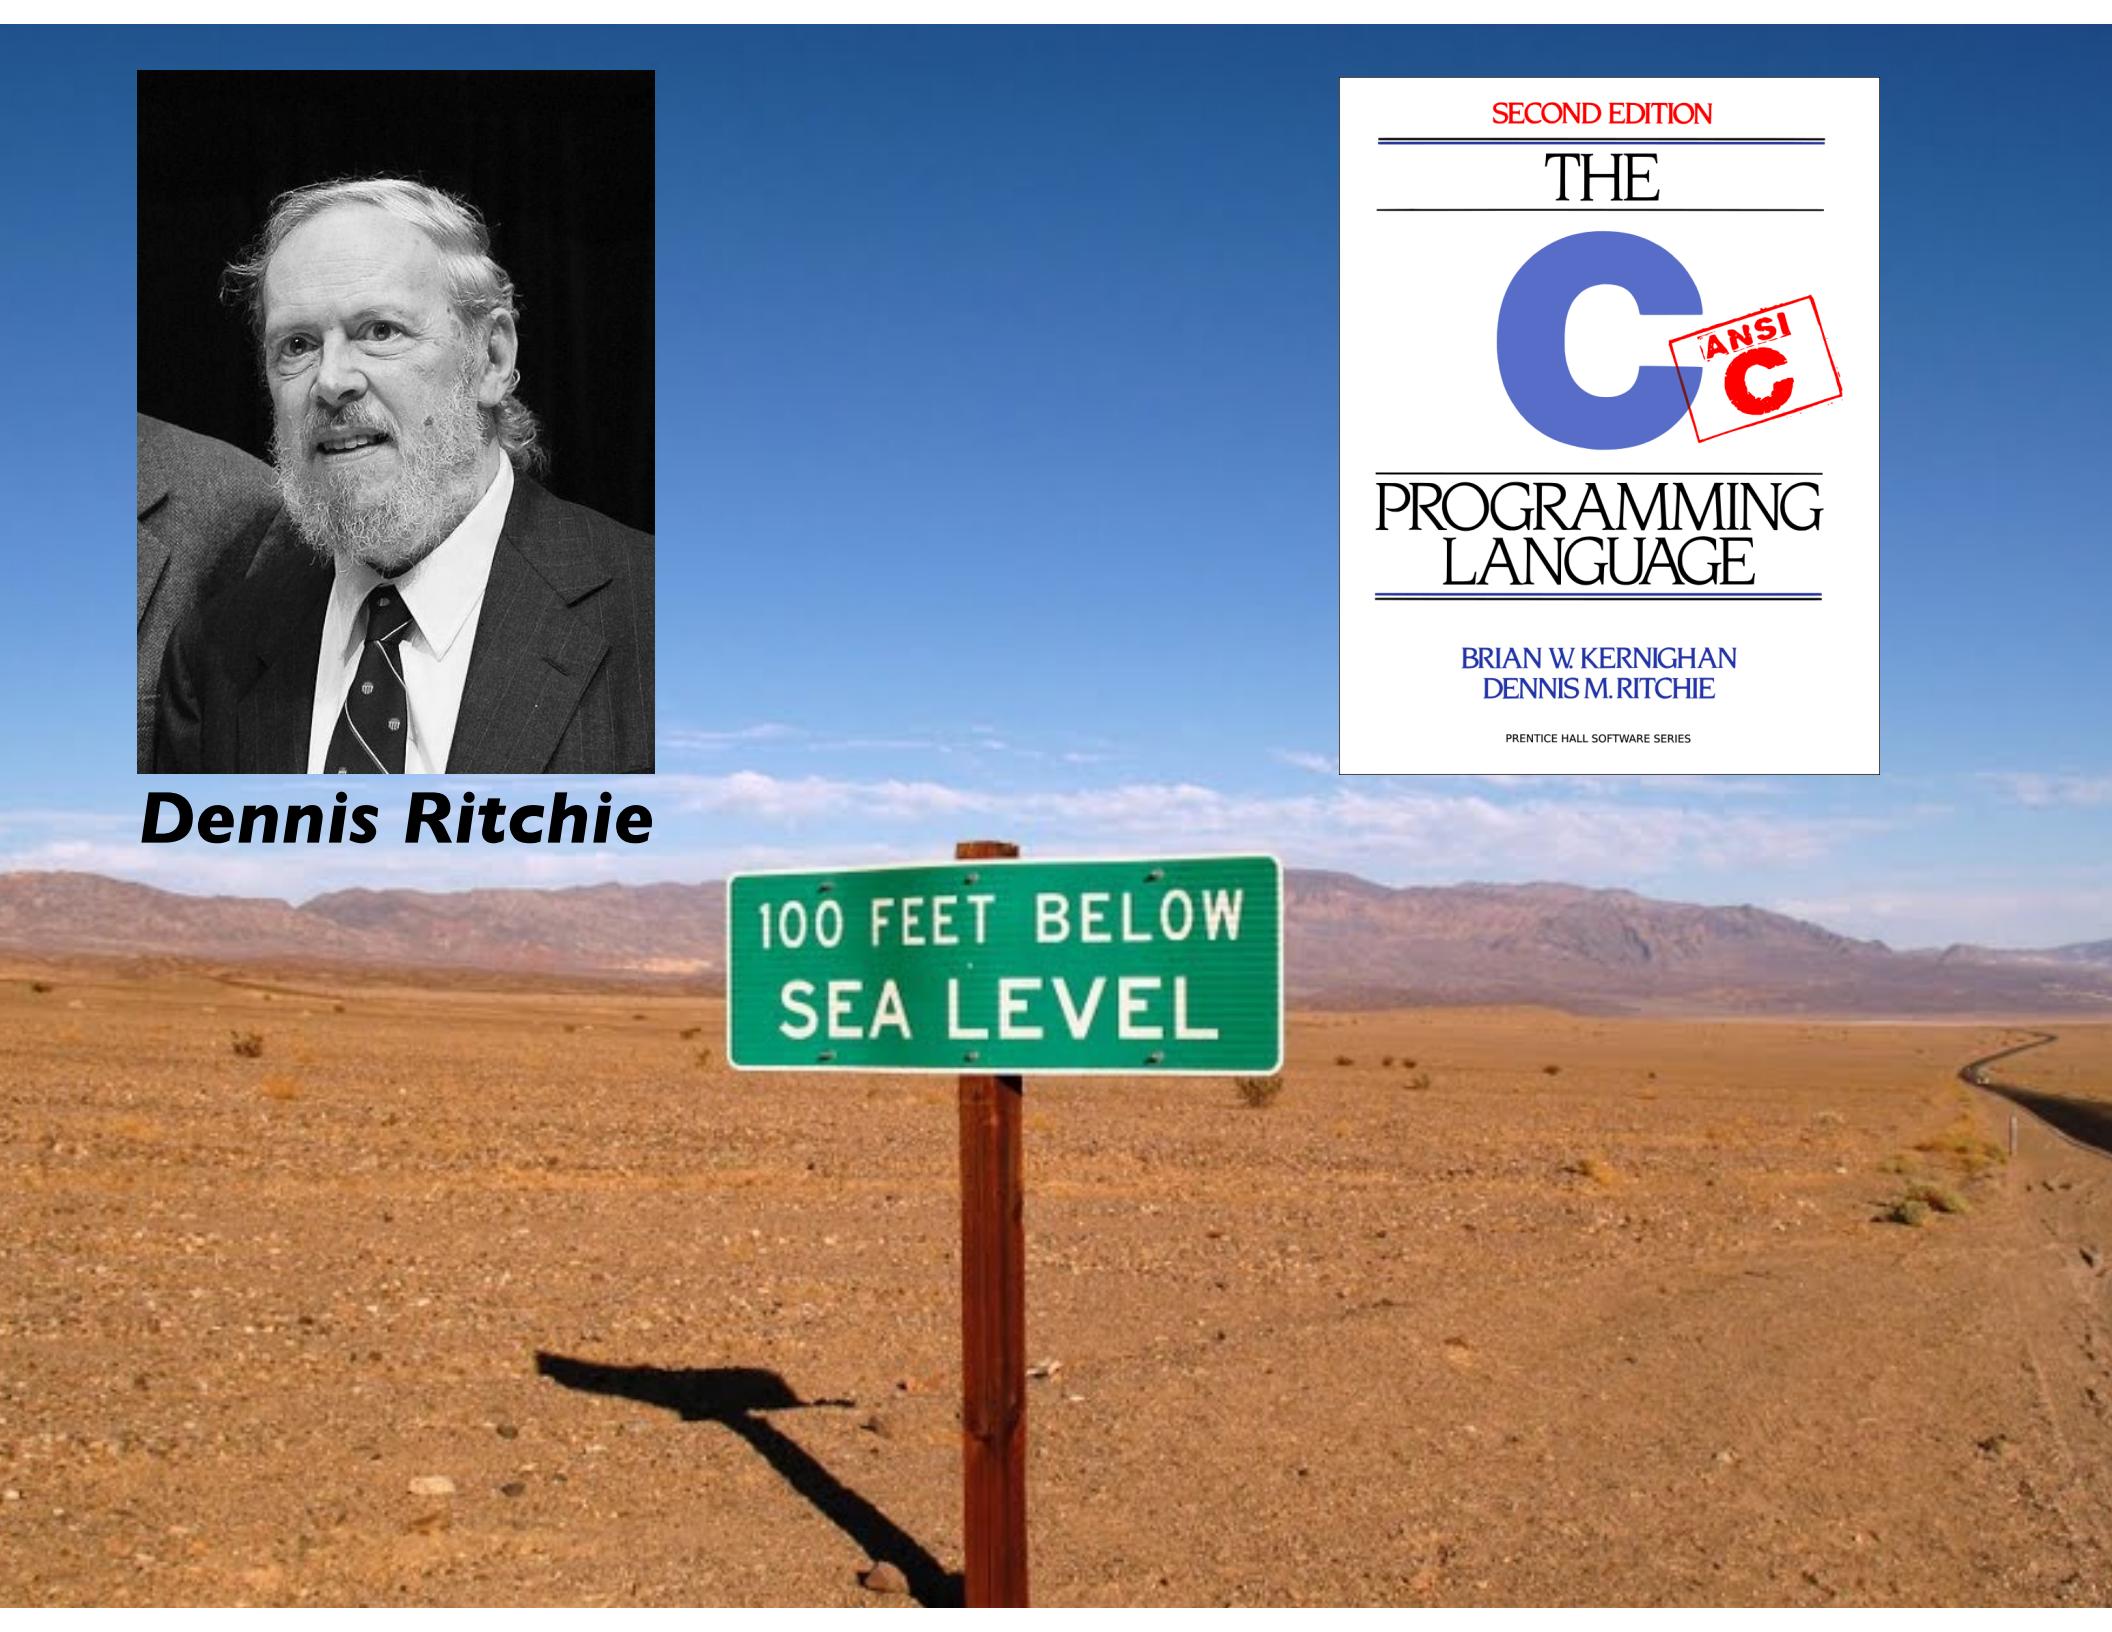
THE



**PROGRAMMING  
LANGUAGE**

BRIAN W. KERNIGHAN  
DENNIS M. RITCHIE

PRENTICE HALL SOFTWARE SERIES

A green road sign on a wooden post in a desert landscape. The sign reads "100 FEET BELOW SEA LEVEL".

100 FEET BELOW  
SEA LEVEL

# C is the language of choice for systems programmers



*Ken Thompson built UNIX using C*

This is not coincidence!  
C features closely model the ISA:  
data types, arithmetic/logical  
operators, control flow, access to  
memory, ... all provided in form  
of portable abstractions

“BCPL, B, and C family of languages are particularly oriented towards system programming, are small and compactly described, and are amenable to translation by simple compilers. They are “close to the machine” in that the abstractions they introduce are readily grounded in the concrete data types and operations supplied by conventional computers, and they rely on library routines for input-output and other interactions with an operating system. ... At the same time, their abstractions lie at a sufficiently high level that, with care, portability between machines can be achieved.”

— Dennis Ritchie

# The C Programming Language

“C is quirky, flawed, and an enormous success”

— Dennis Ritchie

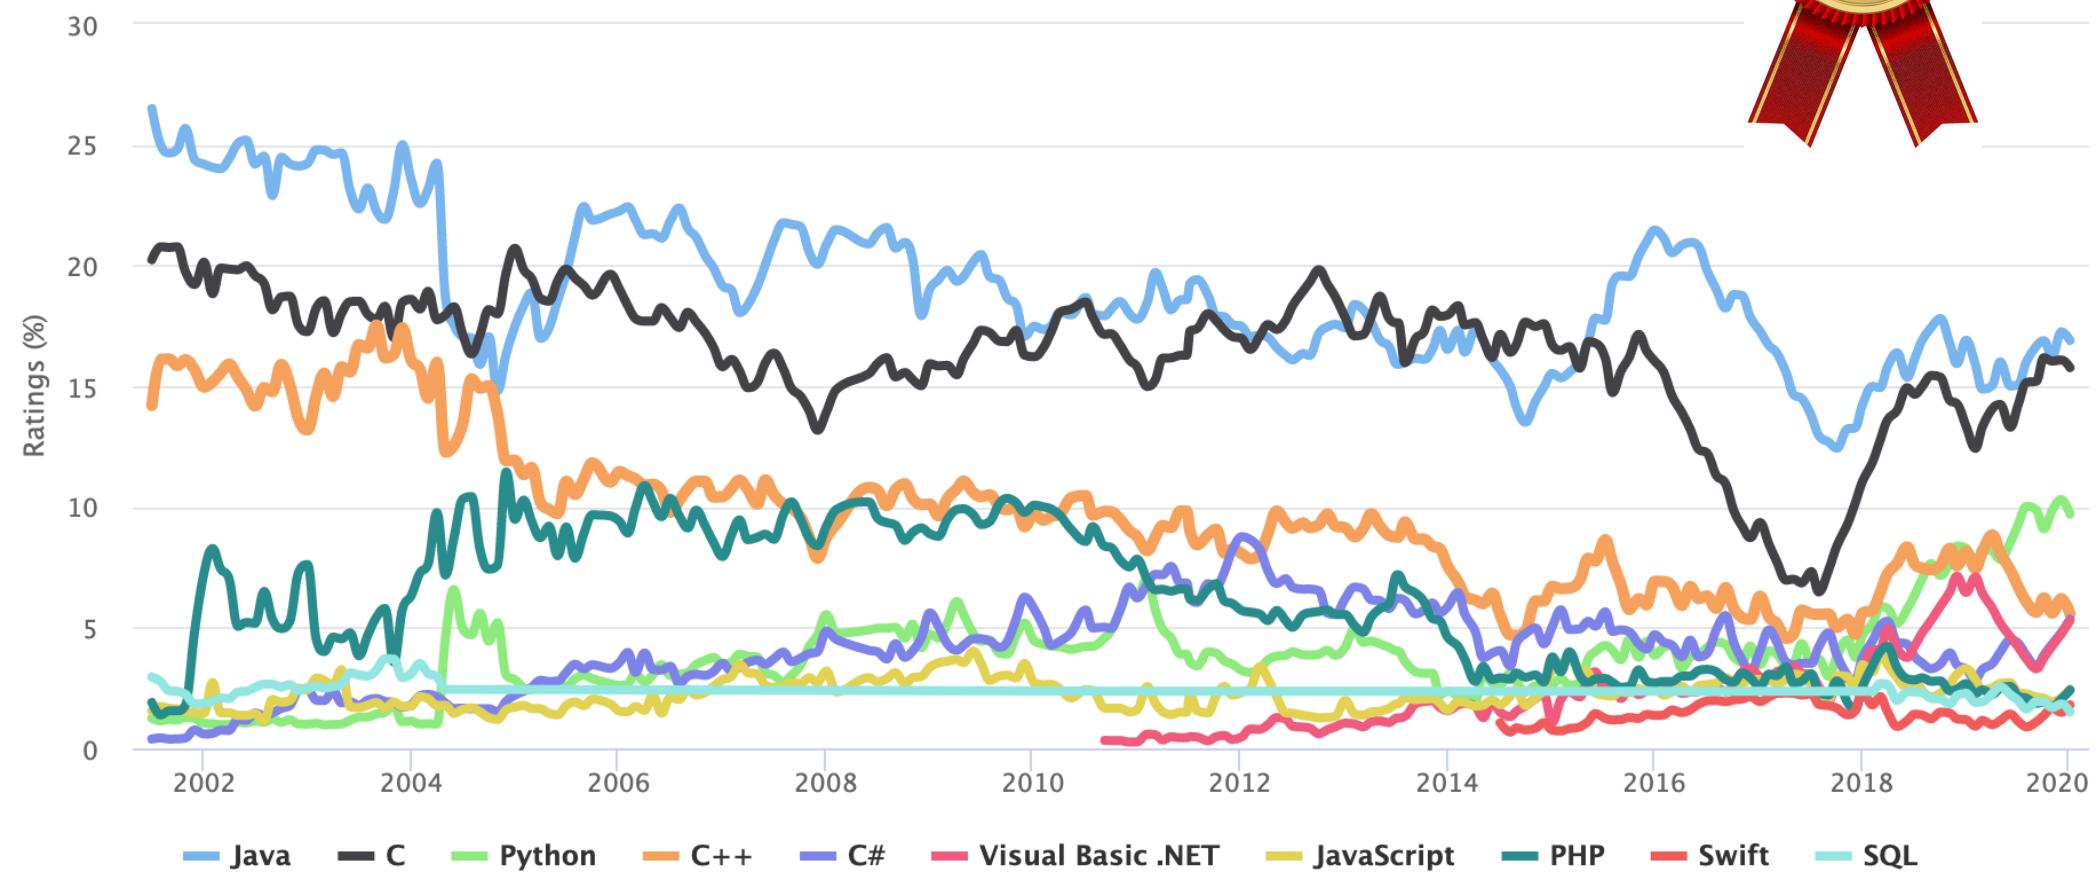
“C gives the programmer what the programmer wants; few restrictions, few complaints”

— Herbert Schildt

“C: A language that combines all the elegance and power of assembly language with all the readability and maintainability of assembly language”

— Unknown

January 2020 headline:  
Programming Language of the Year 2019...



Programming language popularity over time

# Compiler Explorer

is a neat interactive tool to see translation from C to assembly.  
Let's try it now!

The screenshot shows the Compiler Explorer interface. On the left, the 'C source #1' tab contains the following C code:

```
1 int global = 107;
2
3 void main(void)
4 {
5     global = global + 1;
6 }
```

The code editor has a red box around the dropdown menu above it, which is set to 'C'. To the right, the 'ARM gcc 5.4.1 (none) (Editor #1, Compiler #1)' tab displays the generated assembly code:

```
1 main:
2     ldr    r2, .L2
3     ldr    r3, [r2]
4     add    r3, r3, #1
5     str    r3, [r2]
```

The assembly tab also has three red boxes highlighting the compiler settings: 'ARM gcc 5.4.1 (none)' (with a checked green checkbox), and '-Og'.

<https://godbolt.org>

Configure settings to follow along:

C

ARM gcc 5.4.1(none)

-Og

# Why C?

## **Higher-level abstractions, structured programming**

Named variables, constants

Arithmetic/logical operators

Control flow

## **Portable**

Not tied to particular ISA or architecture

## **Low-level enough to get to machine when needed**

Bitwise operations

Direct access to memory

Embedded assembly, too!

# Know your tools!

## **Assembler as**

Transform assembly code (text)  
into object code (binary machine instructions)  
Mechanical rewrite, few surprises

## **Compiler gcc**

Transform C code (text)  
into object code  
(likely staged C → asm → object)  
Complex translation, high artistry

# Make

Build/compile your code using **make**

A **Makefile** describes how to build a piece of software  
Rules, dependencies, and recipes

```
mine.bin: mine.s
    arm-none-eabi-as mine.s -o mine.o
    arm-none-eabi-objcopy mine.o -O binary mine.bin

install: mine.bin
    rpi-install.py mine.bin

clean:
    rm *.o *.bin
```

# Make

Build/compile your code using **make**

A **Makefile** describes how to build a piece of software  
Rules, targets, dependencies, and recipes

blink.bin: blink.s

Rule

```
arm-none-eabi-as blink.s -o blink.o
```

```
arm-none-eabi-objcopy blink.o -O binary blink.bin
```

install: blink.bin

Dependency

```
rpi-install.py blink.bin
```

Recipe

Target

clean:

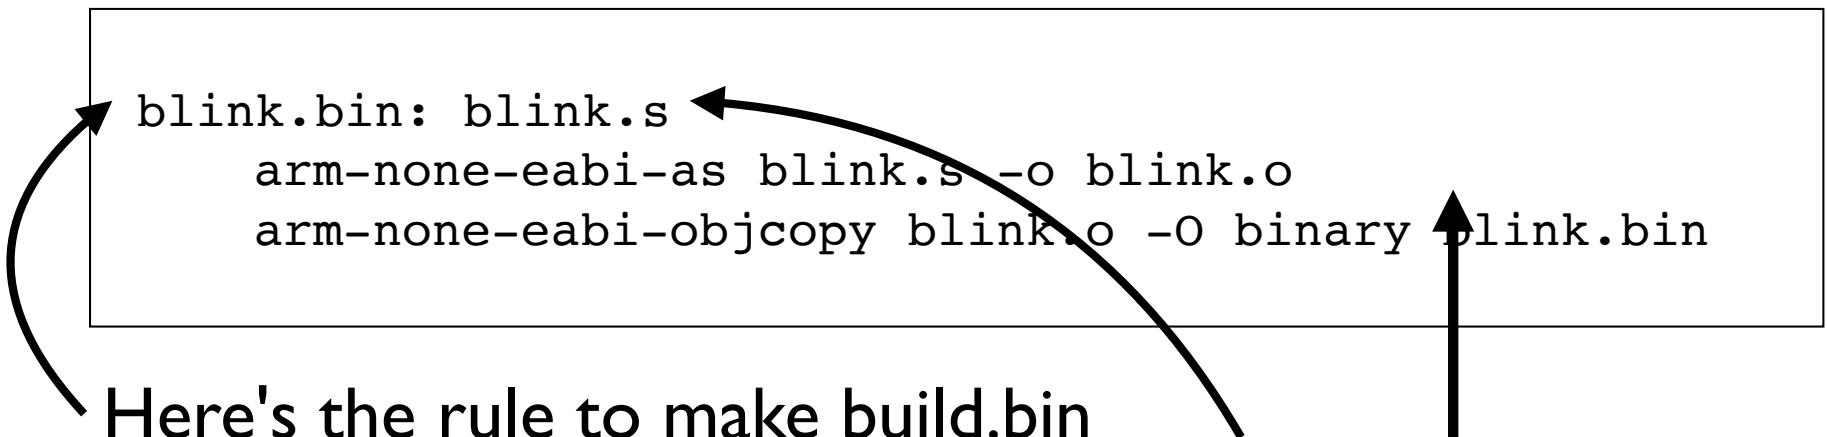
```
rm *.o *.bin
```

# Make

Build/compile your code using **make**

A **Makefile** describes how to build a piece of software

Rules, targets, dependencies, and recipes



Here's the rule to make build.bin

It depends on build.s; if that changes, run this recipe again

Here are the steps (recipe) to produce build.bin

# Make

Build/compile your code using **make**

A **Makefile** describes how to build a piece of software  
Rules, targets, dependencies, and recipes

```
blink.bin: blink.s
    arm-none-eabi-as blink.s -o blink.o
    arm-none-eabi-objcopy blink.o -O binary -o blink.bin
```

Here's the rule to make build.bin

It depends on build.s; if that changes, run this recipe again

Writing out all recipes explicitly like this is onerous, so make has  
all kinds of ways to match patterns, define variables, etc.

# Make pattern rules

**NAME = myprogram**

**CFLAGS = -Og -g -Wall -std=c99 -ffreestanding**

**LDFLAGS = -nostdlib -e main**

**all : \$(NAME).bin**

**% .bin: %.elf**

**arm-none-eabi-objcopy \$< -O binary \$@**

**% .elf: %.o**

**arm-none-eabi-gcc \$(LDFLAGS) \$< -o \$@**

**% .o: %.c**

**arm-none-eabi-gcc \$(CFLAGS) -c \$< -o \$@**

**% .o: %.s**

**arm-none-eabi-as \$< -o \$@**

# Bare-metal vs. Hosted

The default build process for C assumes a *hosted* environment.

What does a hosted system have that we don't?

- standard libraries
- standard start-up sequence
- OS services

To build bare-metal, our Makefile disables these defaults

We supply our own replacements where needed

# Build settings for bare-metal

Compile freestanding

**CFLAGS = -ffreestanding**

Link excludes standard library and start files

**LDFLAGS = -nostdlib**

Link with gcc if need division (b/c no ARM divide instruction)

**LDLIBS = -lgcc**

Write our own code for all libs and start files

This puts us in an exclusive club...

```
int getRandomNumber()
{
    return 4; // chosen by fair dice roll.
              // guaranteed to be random.
}
```