

Admin

- Assign 2
Functionality test results out soon
Revise and resubmit on open issues will be avail
- Assign 3
printf perseverance and pride!
- Lab4
stack, heap, build process

Today: Thanks for the memory!

Linker script `memmap.ld`

Loading, how an executable file becomes a running program

Address space layout, organization of memory

Runtime stack, stack frames, stack allocation

Heap allocation, malloc and free


Linking

Combines one or more object files into executable

Resolve inter-module references

Consolidate code, data sections across all modules

Arrange sections in output file in proper order



```
0000000040000000 T _start start.o
0000000040000008 t hang start.o
000000004000000c t blink clock.o
0000000040000050 T main clock.o
0000000040000150 T _cstart cstart.o
00000000400001a8 T gpio_init gpio.o
00000000400001ac T gpio_set_input gpio.o
00000000400001b0 T gpio_set_output gpio.o
00000000400001b4 T gpio_set_function gpio.o
00000000400001b8 T gpio_get_function gpio.o
00000000400001c0 T gpio_write gpio.o
00000000400001c4 T gpio_read gpio.o
00000000400001cc T timer_get_ticks timer.o
00000000400001d4 T timer_init timer.o
00000000400001d8 T timer_delay_us timer.o
0000000040000210 T timer_delay timer.o
0000000040000234 T timer_delay_ms timer.o
0000000040000288 R __bss_end
0000000040000288 R __bss_start
```

`clock.o timer.o gpio.o start.o cstart.o -> clock.elf`

memmap.ld (linker script)

```
SECTIONS
{
    .text 0x40000000 : { *(.text.start) *(.text*) }
    .rodata : { *(.rodata*) }
    .data : { *(.sdata*) }
    __bss_start = .;
    .bss : { *(.bss*) }
    __bss_end = .;
}
```

What is special about `.text.start` and why must it go first?

What is the significance of `0x40000000`?

Program start sequence

How is a program loaded into memory?

What must happens to start executing a program?

What is known about state of registers or contents of memory?

Does execution actually start at `main()`

👁️ **Read our files!** 👁️

[\\$CS107E/src/start.s](#)

[\\$CS107E/src/cstart.c](#)

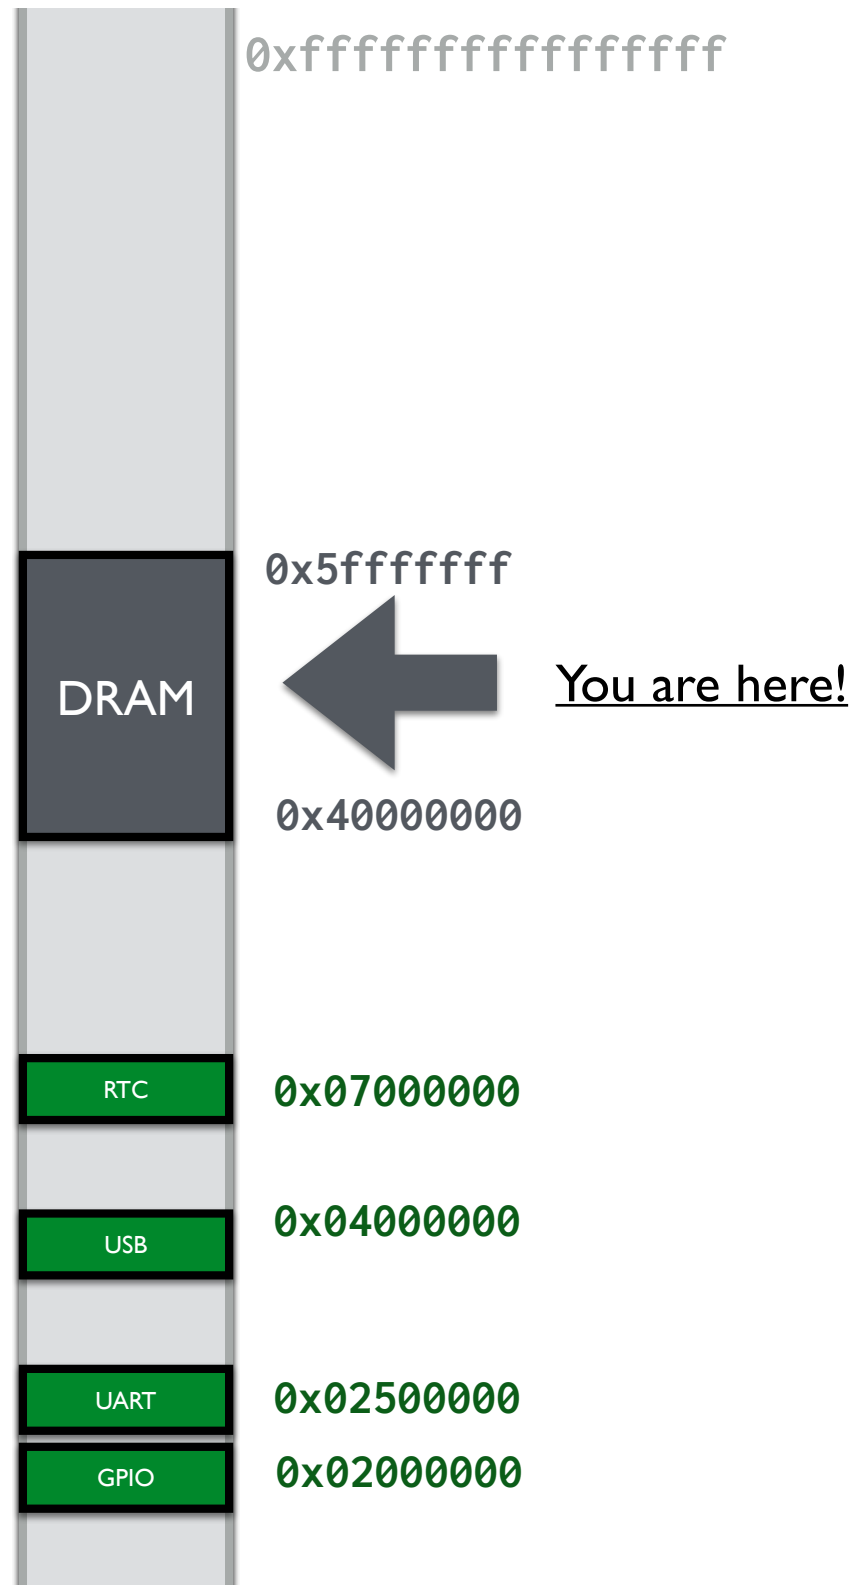
[\\$CS107E/lib/memmap.ld](#)

Memory map

64-bit address space

0x0 - 0xffffffffffffffff

512MB physical RAM at
0x40000000-0x5fffffffff



SECTIONS

```
{
    .text 0x40000000 :{ *(<.text.start)
                        *(<.text*)}
    .rodata :          { *(<.rodata*) }
    .data :            { *(<.data*) }
    __bss_start      = .;
    .bss :             { *(<.bss*) }
    __bss_end        = . ;
}
```

(zeroed data) .bss


(initialized data) .data

(read-only data) .rodata

.text

```
$ xfel write 0x40000000 clock.bin
```

```
$ xfel exec 0x40000000
```



0x50000000

_cstart

main

_start:

lui sp,0x50000

jal _cstart

```
void _cstart(void) {
    char *bss = &__bss_start__;
    while (bss < &__bss_end__)
        *bss++ = 0;
}
main();
}
```

__bss_end

__bss_start

clock.bin

0x40000000

Let's use gdb to watch in action

```
Breakpoint 2, sqr (v=v@entry=3) at program.c:5
5      int sqr(int v) {
(gdb) bt
#0  sqr (v=v@entry=3) at program.c:5
#1  0x0000000040000094 in delta (a=a@entry=3, b=b@entry=7) at program.c:10
#2  0x00000000400000d8 in main () at program.c:15
#3  0x0000000040000048 in _cstart () at cstart.c:7
#4  0x0000000040000010 in _start_gdb ()
(gdb) step
6      return v * v;
(gdb) p v
$1 = 3
(gdb) |
```

Some useful gdb commands to learn

run, step, next, continue

print, x, display

list, disassemble, info registers

breakpoint

backtrace, info frame, up/down

Also repeat, tab completion, history

 **Read our course guide** 


<http://cs107e.github.io/guides/gdb/>

```
// start.s
lui    sp,0x5000
jal    ra,main
```

```
void main(void)
{
    delta(3, 7);
}
```

```
int delta(int a, int b)
{
    int diff = sqr(a) - sqr(b);
    return diff;
}
```

```
int sqr(int v)
{
    return v * v;
}
```



Memory diagram not to scale :-)

Frame pointer

Designate register `s0` for use as **frame pointer** `fp`

Stack pointer `sp` is top of stack, additional register `fp` marks boundary of current stack frame

Each frame saves previous `fp` — this gives reliable way to backtrace entire stack

CFLAGS to enable: `-f-no-omit-frame-pointer`

Add instructions to prolog/epilog that set up `fp` on entry and restore saved on exit

Tracing stack frames

Prolog

Adjust stack pointer to make space ($16 + \text{rest}$)

Store ra and fp to stack

Set **fp** to where sp was (end of prev stack frame)

Body

fp stays anchored


access data on stack **fp**-relative

fp offsets remain stable (even if sp changes)

Epilog

Restore ra, fp to saved values


Adjust sp to remove frame



Linked chain of frame pointers

Can traverse from innermost frame (**sqr**) to frame of its caller (**delta**), from there to its caller (**main**) ...

```
// start.s  
  
// init fp = 0 as termination  
li fp,0  
lui sp,0x50000  
jal _cstart
```



other =
more saved regs, locals, scratch

*Deep dive into full frame
coming up in this week's lab!*

Frame pointer tradeoffs

- + Anchored fp, offsets are constant
- + Standard frame layout enables runtime introspection
- + Backtrace for debugging, instrumentation
- fp register removed from general pool
- Add'l 3 instructions in prolog,epilog to manage fp
- Adds 8 bytes to frame size, another saved register

Aside: "Fedora's tempest in a stack frame"

<https://developers.redhat.com/articles/2023/07/31/frame-pointers-untangling-unwinding>

<https://fedoraproject.org/wiki/Changes/fno-omit-frame-pointer>

<https://lwn.net/Articles/919940/>

We have global storage ...

- + **Convenient**

 - Fixed location, shared across entire program

 - No explicit allocate/deallocate

- + **Fairly efficient, plentiful**

 - (But cost to send over serial line to bootloader)

- +/- **Scope and lifetime is global**

 - No encapsulation, hard to track use/dependencies

 - One shared namespace, possibility of conflicts

 - Frowned upon stylistically

... and we have stack storage ...

- + **Convenient**

 - Automatic alloc/dealloc on function entry/exit

- + **Efficient, fairly plentiful**

 - (But finite size limit on total stack usage)

- +/- **Scope/lifetime dictated by control flow**

 - Private to stack frame

 - Does not persist after function exits

Why do we also need a heap?

An example:

`code/heap/names.c`

Dynamic storage

- + **Programmer controls scope/lifetime**

 - Versatile, precise

 - Works for situations where global/stack do not

- **Needs software runtime support**

 - Library module to manage heap memory

 - Functions to allocate/deallocate memory (explicit calls by client)

- **C version is low on safety**

 - No type safety (raw `void*`, size in number of bytes)

 - Much opportunity for error

 - (allocate wrong size, use after free, double free)

Heap module interface

```
void *malloc (size_t nbytes);  
void free (void *ptr);
```

What is a void* pointer?

"Generic" pointer, holds a memory address

Type of pointee is not specified, could be any type of data

What you can do with a void*

Pass to/from function, assignment

What you cannot do with a void*

No dereference without cast

No pointer arithmetic without cast (scaling unknown!)

No array indexing (size of pointee unknown!)

How to implement a heap



Drawing by Jane Lange

Bump allocator

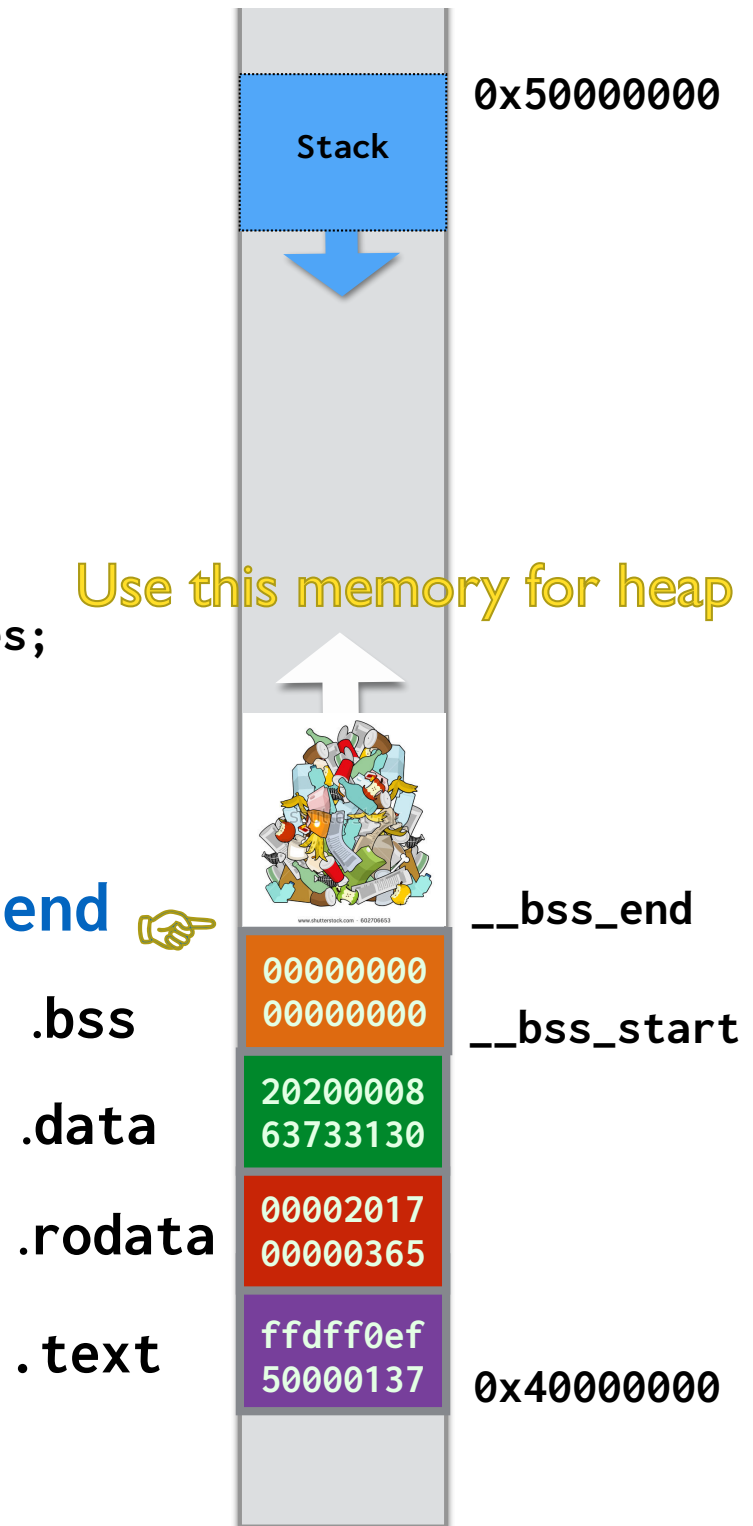
```
void *sbrk(int nbytes)
{
    static void *_cur_heap_end = &__bss_end;

    void *prev_end = cur_heap_end;
    cur_heap_end = (char *)cur_heap_end + nbytes;
    return prev_end;
}
```


void *cur_heap_end 

- Keep pointer to end of heap segment
- Service malloc request adjusts pointer upward
- Every request extends/grows heap segment
- No reuse/recycle

Use this memory for heap



Tracing the bump allocator



Bump Memory Allocator

`code/heap/malloc.c`

Evaluate bump allocator

- + Operations super-fast
- + Very simple code, easy to verify, test, debug
- No recycling/re-use
 - (in what situations will this be problematic?)
- Sad consequences if `sbrk()` advances into stack
 - (what can we do about that?)


Pre-block header, implicit list

Each square represents 8 bytes, header records size of payload in bytes


`p1 = malloc(32)`




`p2 = malloc(40)`




`p3 = malloc(48)`



`free(p2)`



`p4 = malloc(16)`



Header struct on each block

```
struct header {
    unsigned int size;
    unsigned int status;
};                                     // sizeof(struct header) = 8 bytes

enum { IN_USE = 0, FREE = 1 };

void *malloc(size_t nbytes)
{
    nbytes = roundup(nbytes, 8);
    size_t total_bytes = nbytes + sizeof(struct header);

    struct header *hdr = sbrk(total_bytes); // extend end of heap
    hdr->size = nbytes;
    hdr->status = IN_USE;
    return hdr + 1;      // return address at start of payload
}
```


Challenge for malloc client

Correct allocation (size in bytes)

Correct access to block (within bounds, not freed)

Correct free (once and only once, at correct time)

What happens if you...

- forget to free a block after you are done using it?
- access a memory block after you freed it?
- free a block twice?
- free a pointer you didn't malloc?
- access outside the bounds of a heap-allocated block?

Challenge for malloc implementor

just malloc is easy 😎

malloc with free is hard 🤔

Efficient malloc with freeYikes! 😓

Complex code (pointer math, typecasts)

Critical system component

correctness is non-negotiable!

Thorough testing is essential

Survival strategies:

draw lots of pictures

printf (you've earned it!!)

early tests on inputs small enough to trace by hand if need be

build up to more complex tests